

VIKASH JAISWAL

Lead Platform Engineer · Principal DevOps Engineer · AI Systems Architect · SRE

📍 Noida / Gurgaon / Delhi NCR, India | 📞 +91-8588800287 | ✉️ vikashjaiswal.486@gmail.com

Available · 15-Day Notice · Open to Remote / PAN India / Global Roles

Professional Summary

Lead Platform Engineer and Principal DevOps Architect with **15+ years** of progressive experience spanning **Broadcast Operations** → **OTT/Media Engineering** → **Multi-Cloud Infrastructure** → **Site Reliability Engineering** → **Platform Engineering** → **Enterprise GenAI Systems**. Proven track record designing, operating, and modernizing large-scale production environments — from 450+ live television channels at Tata Sky to 19 Kubernetes clusters across 7 global regions at Devo Technology.

Builder of **AI FORGE**, an enterprise AI Agent Operating System on AWS Bedrock that reduced engineer troubleshooting ramp-up by ~60%. Author of **forge-router**, a production-ready multi-LLM gateway routing across 8 AI providers. Certified AWS Solutions Architect with deep expertise in Terraform (including **Custom Provider development via Plugin SDK**), GitOps, DevSecOps, observability engineering, and AI platform design. Currently pursuing LLMops, Agentic AI, RAG architectures, and AI Gateway design as the next platform frontier.

Impact at a Glance

Metric	Achievement
Experience	15+ years across Broadcast, OTT, Cloud, SRE, Platform Engineering, GenAI
Scale	19+ Kubernetes clusters · 7 global regions · 450+ live TV channels
AI Innovation	Built AI FORGE (enterprise GenAI OS) + forge-router (multi-LLM gateway, 8 providers)
Automation	EBS storage incidents: 45-min manual → 3-min zero-touch Lambda automation
Efficiency	~60% reduction in engineer troubleshooting ramp-up via AI-assisted operations
IaC Depth	Custom Terraform Provider development using Plugin SDK (rare, senior-level skill)
Security	Zero-secret architecture: Vault/OpenBao across all K8s, Jenkins, GitLab pipelines
Observability	GitOps-managed Prometheus + OpenTelemetry + Grafana across multi-cloud environments
Multi-cloud	AWS ★★★★★ · GCP ★★★★★ · Azure ★★★★★ · IBM Cloud ★★★★★

Career Progression

timeline

title 15+ Years of Engineering Evolution	
2009–2010	: TVS-E / Nokia : Service Engineer : Hardware Diagnostics & Mobile Platform
2010–2012	: Essel Shyam (Star · Disney · Hotstar) : Broadcast Engineer : Master Control Room · Playout Automation
2012–2019	: Tata Sky Ltd : Assistant Manager – Broadcast Operations : 450+ TV Channels · OTT · AWS Lambda · HLS/DASH
2019–2022	: Hughes Communications India : Assistant Manager – Cloud & RF Engineer : AWS Media · VSAT · Cloud Migration
2022–2023	: GlobalXperts Inc : Senior Cloud Consultant : NHL Live Streaming · AWS MediaLive · Terraform
2023–Now	: Devo Technology : Lead DevOps Engineer · SRE : AI FORGE · 19 K8s Clusters · 7 Regions · AWS Bedrock

Core Competencies

mindmap

root((Vikash Jaiswal))	
Cloud & Infrastructure	
AWS	★★★★★
GCP	★★★★
Azure	★★★★
IBM Cloud	★★★★
Multi-cloud Architecture	
FinOps	
Infrastructure as Code	
Terraform	★★★★★
Terragrunt	★★★★★
Custom Provider SDK	
Ansible	★★★★
Policy as Code	
Containers & Platform	
Kubernetes	★★★★★
Docker	★★★★★
ArgoCD	★★★★★
Helm	★★★★

- OpenShift
- GitOps
- CI/CD
 - Jenkins ★★★★★
 - GitHub Actions ★★★★★
 - GitLab CI ★★★★★
 - Azure DevOps
- Observability & SRE
 - Prometheus
 - Grafana
 - OpenTelemetry
 - Dynatrace
 - Datadog
 - SLI/SLO/Error Budgets
 - DORA Metrics
- AI & GenAI
 - Claude Code
 - AWS Bedrock
 - Gemini Enterprise
 - Cursor AI
 - RAG Architectures
 - LLMOps
 - Multi-Agent Systems
 - MCP
 - AI Gateway Design
 - Vector Databases
- Security
 - Vault/OpenBao
 - DevSecOps
 - SAST/SCA
 - Zero-Trust
 - IAM/RBAC

Technical Skills Matrix

Cloud Platforms & AWS Services

Category	Technologies
Cloud Platforms	AWS · Google Cloud (GCP) · Microsoft Azure · IBM Cloud
Compute & Networking	EC2 · EKS · Lambda · VPC · Route53 · CloudFront · API Gateway · ELB
Storage & Data	S3 · RDS · DynamoDB · EBS · ElastiCache · Elasticsearch
Automation & Events	Lambda · Step Functions · SNS · SQS · EventBridge · CloudWatch
Security & Identity	IAM · SSO · Secrets Manager · KMS · SSM · Organizations
AI & ML	AWS Bedrock · Amazon Titan · SageMaker · MediaLive · MediaConnect · MediaPackage
GCP Services	GKE · Cloud Run · Cloud Build · Gemini API · Pub/Sub
Azure Services	AKS · Azure DevOps · Azure Monitor · Azure Storage · Virtual Machines

Infrastructure as Code

Technology	Proficiency	Notable Capability
Terraform	★★★★★	Custom Provider development using Plugin SDK + Framework lifecycle
Terragrunt	★★★★★	Multi-environment orchestration, DRY configurations, remote state
Ansible	★★★★	Playbooks, Roles, Jinja2, Inventory management, Configuration drift
Policy as Code	★★★★	OPA, Sentinel, infrastructure compliance automation
Helm	★★★★	Chart authoring, values management, versioned releases

Containers & Orchestration

Technology	Proficiency	Context
Kubernetes	★★★★★	19+ production clusters, 7 global regions, multi-tenancy
Docker	★★★★★	Image optimization, multi-stage builds, registries
ArgoCD	★★★★★	GitOps delivery, Blue-Green, Canary, Rolling deployments
Red Hat OpenShift	★★★★	Enterprise Kubernetes, SDN, RBAC
Amazon EKS	★★★★★	Multi-cluster management, node groups, add-ons

CI/CD Platforms

Platform	Proficiency	Depth
Jenkins	★★★★★	Groovy shared pipeline libraries, master/agent, multi-branch
GitLab CI/CD	★★★★	Multi-stage pipelines, environments, protected branches
GitHub Actions	★★★★	Reusable workflows, OIDC auth, matrix builds
Azure DevOps	★★★★	Pipelines, release gates, artifact management
ArgoCD (GitOps)	★★★★★	App of Apps, ApplicationSets, Sync waves

Observability & SRE Stack

Tool	Role
Prometheus	Metrics collection, PromQL, alerting rules, federation
Grafana	Dashboards, alert routing, GitOps-managed configuration
OpenTelemetry	Collectors, DaemonSet/Sidecar patterns, traces, spans
Dynatrace	Full-stack monitoring, APM, infrastructure intelligence
Datadog	Cloud monitoring, APM, log management, streaming metrics
ELK Stack	Elasticsearch, Logstash, Kibana — log aggregation and analysis
Loki	Log aggregation, label-based querying, Grafana integration
Alertmanager	Alert deduplication, routing, silencing, escalation

AI & GenAI Engineering Stack

Category	Technologies & Concepts
AI Platforms	Claude Code · AWS Bedrock · Gemini Enterprise · Cursor AI · GitHub Copilot · Codex CLI
Agentic AI	Multi-Agent Systems · Agentic Workflows · AI FORGE (own project)
RAG Architectures	Chunking · Embeddings · Hybrid Search (BM25 + Semantic) · Cross-encoder Reranking

Category

Technologies & Concepts

Vector Databases	OpenSearch Vector DB · Qdrant · ChromaDB · HNSW Indexing · pgvector
LLM Frameworks	LangChain · LangGraph · LangSmith · Prompt Engineering
AI Observability	Langfuse · AI Telemetry · Prompt Monitoring · LLM-as-a-Judge
AI Infrastructure	AI Gateway Design · Multi-LLM Routing · forge-router (own project)
AI Governance	PII Redaction · Guardrails · AuthN/AuthZ · Deny-list Safety
Protocols	Model Context Protocol (MCP) · MCP Tools/Resources/Prompts

Security & DevSecOps

Category

Technologies

Secrets Management	HashiCorp Vault · OpenBao · Dynamic Credentials · Token Lifecycle
Pipeline Security	SAST · SCA · Container Image Scanning · Secret Scanning · IaC Policy
Identity & Access	IAM · RBAC · Zero-Trust Architecture · mTLS · Least Privilege
Compliance	Compliance as Code · SBOM Generation · Artifact Signing · Audit Trails
Runtime Security	K8s Network Policies · Pod Security · Anomaly Detection

Programming & Scripting

Language

Proficiency

Primary Use

Python	★★★★★	Automation, Lambda, Boto3, API scripting, AI tooling
Bash / Shell	★★★★★	Infrastructure scripts, wrapper automation, CI tooling
Groovy	★★★★★	Jenkins shared libraries, pipeline DSL
YAML / JSON	★★★★★	K8s manifests, Terraform configs, pipeline definitions
TypeScript	★★★	Next.js portfolio, frontend engineering
SQL	★★★★★	MySQL queries, RDS administration

Databases & Storage

Technology

Context

MySQL / Amazon RDS	Multi-region schemas, production DBA operations at Devo
Elasticsearch / OpenSearch	Log analytics, vector search, HNSW indexing
Amazon DynamoDB	Serverless event-driven architectures, high-throughput
Amazon S3	Object storage, data lake patterns, event triggers
Vector Databases	OpenSearch, Qdrant, ChromaDB for RAG pipelines

Media, OTT & Broadcast Domain

Category

Technologies & Experience

AWS Media Services	MediaLive · MediaConnect · MediaPackage · CloudFront
Streaming Protocols	HLS · DASH · Adaptive Bitrate (ABR) · Manifest generation
Broadcast Systems	Harris Automation · Playout Systems · Master Control Room
Satellite/RF	VSAT · Link Budget Calculations · EIRP · Transponder Bandwidth
CDN Operations	Akamai · AWS CloudFront · Edge distribution · Origin management
Video Processing	FFmpeg · Transcoding · Content packaging · DRM
OTT Platforms	Tata Sky Binge · Subscriber delivery · Content onboarding · Metadata

Professional Experience

Devo Technology — Noida, India

Lead DevOps Engineer / Site Reliability Engineer | Apr 2023 – Present

Devo Technology is a cloud-native security analytics and SIEM platform serving global enterprise customers. Led platform engineering, infrastructure automation, observability modernization, and AI-assisted operational intelligence across 19+ Kubernetes clusters in 7 global regions.

AI-Assisted Operations (AI FORGE):

- Designed and delivered **AI FORGE** — enterprise GenAI operational platform on AWS Bedrock integrating Claude AI, reducing engineer troubleshooting ramp-up by **~60%**
- Built 11 domain-specific AI skill modules covering Kubernetes, Jenkins, GitLab, Vault, Jira, Confluence, MySQL, observability, LINQ query engine, and multi-region operations
- Implemented persistent cross-session memory architecture enabling context retention, incident continuity, and knowledge accumulation across engineering shifts
- Created zero-secret wrapper architecture: 15 wrapper scripts injecting credentials at runtime — credentials never appear in AI context, git history, or logs
- Deployed structural safety deny-list for destructive operations (`rm -rf`, `kubectl rollout restart`, `git push origin main`) requiring explicit human confirmation

Infrastructure Engineering:

- Developed **Custom Terraform Provider** using Plugin SDK and Framework lifecycle management — built internal infrastructure abstractions not available in off-the-shelf providers
- Standardized Terraform/Terragrunt IaC across APAC, EU, and US production regions — single configuration change propagated consistently across three regions
- Managed 19+ Amazon EKS clusters across 7 global regions: EU, US, US3, APAC, Santander, GCP/Telefonica, NCSC Bahrain
- Integrated **OpenBao/Vault** across all Kubernetes workloads, Jenkins pipelines, GitLab CI/CD, and multi-region infrastructure — full elimination of hardcoded credentials

Observability & SRE:

- Architected unified observability platform: OpenTelemetry collectors + Prometheus + Grafana with **GitOps-managed alert lifecycle** (all thresholds, routing, and escalation trees as YAML in Git)
- Deployed OpenTelemetry collectors as DaemonSet and Sidecar patterns for full-spectrum telemetry coverage across distributed services
- Built automated log-forwarding via Loki and ELK stacks piping context-rich error traces into SRE operations boards, cutting incident discovery time significantly
- Resolved critical **JVM OutOfMemoryError** incidents on mission-critical CaixaBank banking cluster via heap-dump analysis, object isolation diagnosis, and systemic cache pattern fixes

Automation:

- Built **EBS Storage Expansion Automation Platform**: CloudWatch → SNS → API Gateway → Lambda chain; eliminated 45-minute on-call manual response with 3-minute zero-touch resolution
- Designed self-healing operational workflows where infrastructure health was continuously validated and remediation playbooks automatically executed on service failures
- Developed Ansible Playbooks and reusable Roles for configuration management, deployment automation, and operational standardization across Linux environments

Technology Stack: AWS · GCP · IBM Cloud · Kubernetes · EKS · ArgoCD · Helm · Terraform · Terragrunt · Ansible · Jenkins · GitLab · OpenTelemetry · Prometheus · Grafana · Vault/OpenBao · Python · Claude/Bedrock · MySQL

GlobalXperts Inc — Noida, India

Senior Cloud Consultant | *Feb 2022 – Apr 2023*

Cloud modernization advisory and technical automation for global enterprise clients. Owned production operations for NHL live streaming infrastructure on AWS.

- Owned production operations for **NHL live streaming** on AWS MediaLive, MediaConnect, CloudFront, Route53, IAM under strict SLA requirements
- Architected hybrid networking: custom VPC subnets, load balancers, DNS configurations, CloudFront CDN footprints, and Network Security Groups for high-volume media streams
- Built multi-platform metrics aggregation combining Datadog, CloudWatch, and Grafana for streaming health monitoring across concurrent live broadcast events
- Delivered deep Terraform-based cloud modernization advisory and advanced compute architecture optimization for global enterprise clients
- Led incident response, root cause analysis, and service recovery for customer-facing media services
- Supported cloud-native modernization involving Kubernetes, Terraform, observability platforms, CI/CD pipelines, and platform automation

Technology Stack: AWS · MediaLive · MediaConnect · CloudFront · Terraform · Datadog · CloudWatch · Grafana · Route53 · Python

Hughes Communications India Ltd — Noida, India

Assistant Manager — Cloud & RF Engineer | *Aug 2019 – Feb 2022*

Cloud-integrated streaming media and satellite communications for enterprise and government customers across India.

- Engineered resilient cloud-integrated streaming media architectures for nationwide corporate and **government e-learning platforms** on AWS
- Deployed hybrid communication pathways bridging **AWS Media Services with VSAT satellite networks**, customizing encoder profiles for constrained transport environments
- Conducted advanced **satellite Link Budget calculations**, transponder bandwidth footprint mapping, and EIRP tracking to eliminate carrier degradation at remote edge sites
- Guided enterprise customers through complex infrastructure migration pathways — legacy on-premise to cloud-automated layouts
- Managed VSAT network operations, antenna alignment, RF planning, and performance optimization across satellite communication systems

Technology Stack: AWS · MediaLive · MediaPackage · CloudFront · VSAT · RF Engineering · Python · Terraform

Tata Sky Ltd (Tata Sky Binge OTT) — Delhi NCR, India

Assistant Manager — Broadcast Operations | *Apr 2012 – Jul 2019*

Governed real-time transmission and OTT delivery for one of India's largest DTH operators — 450+ live television channels and the Tata Sky Binge streaming platform serving millions of subscribers.

- Governed live ingestion pipelines and transmission systems for **450+ TV channels** with continuous high-availability SLA obligations
- Orchestrated serverless event-driven automation using **AWS Lambda + Step Functions** for stream transcoding, adaptive packaging (HLS/DASH), and manifest file generation
- Programmed Python health-check modules polling telemetry from bare-metal hardware matrices, hardware transcoders, and central storage arrays
- Partnered with OTT, CDN, product, and software engineering teams supporting **Tata Sky Binge** — content onboarding, metadata workflows, subscriber-facing digital experiences
- Directed major incident management, service continuity planning, and risk mitigation across mission-critical broadcast delivery
- Led broadcast operations: compression systems, RF infrastructure, uplink services, multiplexing, content delivery, and service assurance

Technology Stack: AWS Lambda · Step Functions · HLS/DASH · Akamai CDN · Python · Broadcast Systems · RF Infrastructure

Essel Shyam Communication Ltd — Delhi NCR, India

Broadcast Engineer | *Apr 2010 – Apr 2012*

Broadcast operations for tier-1 media entities including Star India, Disney, Hotstar, and Discovery TV.

- Managed broadcast automation software platforms and master control room distribution for **Star India, Disney, Hotstar, and Discovery TV**
 - Supported Harris automation suites, automated playback loops, and baseband hardware routing to maintain strict SLA uptimes
 - Managed playout systems, monitoring platforms, transmission workflows, content ingestion, and scheduling validation for 24×7 operations
-

TVS-E Service Tec (Nokia) — Delhi NCR, India

Service Engineer | *Aug 2009 – Mar 2010*

- Executed chip-level diagnostics, circuitry path trace verification, and Layer-4 hardware operational evaluations on mobile development platform prototypes
-

Key Projects & Architecture Initiatives

AI FORGE — Enterprise AI Agent Operating System

flowchart TB

```
subgraph INTERFACE["Interface Layer"]
    ENG[Engineer / On-Call SRE]
end
```

```
subgraph INTELLIGENCE["Intelligence Layer – AI FORGE"]
    BRAIN[Global Brain\nCLAUDE.md]
    SKILLS[11 Domain Skills\nBrain Packs]
    MEMORY[Persistent Memory\nfeedback · project · reference · user]
    BRAIN --> SKILLS
    BRAIN --> MEMORY
end
```

```
subgraph EXECUTION["Execution Layer – Wrapper Scripts"]
    MW[maui-wrapper]
    KW[kubectl-wrapper]
    VW[vault-wrapper]
    JW[jenkins-wrapper]
    GW[git-wrapper]
    SW[sql-wrapper]
end
```

```
subgraph PLATFORM["Platform Layer"]
    MAQUI[Maui LINQ Engine\n7 Regions]
    EKS[Amazon EKS\n19 Clusters]
    VAULT[Vault / OpenBao\n5 Regions]
    RDS[MySQL / RDS\nMulti-Region]
    JENKINS[Jenkins CI/CD]
    GITLAB[GitLab]
end
```

```
subgraph GATEWAY["AI Gateway – forge-router"]
    CLAUDE[Claude\npriority 3]
    GROQ[Groq LLaMA\npriority 2]
    GEMINI[Gemini\npriority 1]
    OLLAMA[Ollama Local\npriority 7]
end

subgraph CRED["Security Layer"]
    CRED["~/devo/credentials\nchmod 600 · Single Source"]
    DENY["Deny-List\nrm-rf · reboot · kubectrl rollout · git push main"]
end

ENG -->|Slash Command| SKILLS
SKILLS --> MW & KW & VW & JW & GW & SW
MW --> MAQUI
KW --> EKS
VW --> VAULT
JW --> JENKINS
SW --> RDS

MEMORY -->|Session Context| ENG
INTELLIGENCE -->|Augmented Prompt| GATEWAY
GATEWAY --> CLAUDE & GROQ & GEMINI & OLLAMA

CRED -->|Runtime Injection| MW & KW & VW & JW & GW & SW
DENY -->|Blocks Destructive Ops| EXECUTION
```

Tech Stack: Claude Code · AWS Bedrock · Sonnet 4.6/4.5 · Python · Kubernetes · Vault/OpenBao · Terraform · Jenkins · Prometheus · MCP-aligned Architecture

Impact: ~60% reduction in troubleshooting ramp-up · Zero credential exposure · 11 domain skills · 7 global regions · Cross-session persistent memory

forge-router — Multi-LLM AI Gateway

```
flowchart LR
    PROMPT[Augmented Prompt\nfrom AI FORGE] --> ROUTER

subgraph ROUTER["RouterEngine – forge-router"]
    HC[Health Check\nAll Providers]
    SORT[Priority Sort\n+ Preferred Override]
    FALL[Fallback Chain\nAuto-retry on failure]
    HC --> SORT --> FALL
end

ROUTER --> AG["Antigravity\nGemini 1.5 Flash\nPriority 0"]
ROUTER --> GEM["Gemini CLI\nPriority 1"]
ROUTER --> GROQ["Groq\nLLaMA 3.3 70B\nPriority 2"]
ROUTER --> CLAUDE["Claude\nclaude-3-5-sonnet\nPriority 3"]
ROUTER --> CODEX["Codex\nGitHub Copilot\nPriority 4"]
ROUTER --> OAI["OpenAI\nGPT-4o\nPriority 6"]
ROUTER --> OLLAMA["Ollama\nLocal Models\nPriority 7"]
```

```
style AG fill:#4285F4,color:#fff
style GROQ fill:#FF6B35,color:#fff
style CLAUDE fill:#CC785C,color:#fff
style OLLAMA fill:#2E8B57,color:#fff
```

Tech Stack: Python 3.13 · AsyncIO · Anthropic API · Groq · Gemini · OpenAI · Ollama · Rich TUI · uv

Capabilities: 8 provider adapters · health-check-first routing · automatic fallback · vision/multimodal support · forge doctor diagnostics · interactive TUI

EBS Auto-Remediation — Zero-Touch Storage Automation @ Devo

sequenceDiagram

```
participant CW as CloudWatch
participant SNS as SNS Topic
participant APIGW as API Gateway
participant L as Lambda Function
participant EBS as EBS Volume
participant XFS as XFS Filesystem
participant SLACK as Slack Channel
```

```
Note over CW: Volume utilization > 85%
CW->>SNS: Alarm triggered
SNS->>APIGW: HTTP POST event
APIGW->>L: Invoke handler
L->>L: Validate alarm payload
L->>L: Identify volume + instance
L->>EBS: ModifyVolume API call
EBS-->>L: Volume resize confirmed
L->>XFS: xfs_growfs - online resize
XFS-->>L: Filesystem expanded
L->>SLACK: ✅ Auto-resolved - no action needed
Note over SLACK: Engineers informed AFTER resolution
Note over L,SLACK: Total time: ~3 min | Zero human involvement
```

Before: 30–60 minute on-call manual response — SSH, manual EBS expand, manual XFS resize, ticket
After: 3-minute fully automated resolution — engineers notified of completion, not pulled into execution

Production RAG System — Enterprise Knowledge Platform

flowchart TD

```
subgraph INGEST["Knowledge Ingestion Pipeline"]
    S3[S3 Bucket\nDocs/Runbooks] --> LAMBDA[Lambda\nChunker]
    LAMBDA --> SQS[SQS Queue]
    SQS --> GLUE[Glue ETL\nEmbedding Job]
    GLUE --> EMB["Embedding\nModel"]
    EMB --> TITAN[Amazon Titan]
```

```

    EMB --> OAI_E[OpenAI Embeddings]
    EMB --> BGE[BGE Model]
end

subgraph STORE["Vector Store"]
    OPENSEARCH[OpenSearch\nVector DB\nHNSW Index]
    DDB[DynamoDB\nMetadata Store]
    TITAN & OAI_E & BGE --> OPENSEARCH
    GLUE --> DDB
end

subgraph RETRIEVAL["Retrieval & Ranking"]
    QUERY[User Query] --> BM25[BM25\nKeyword Search]
    QUERY --> SEM[Semantic\nVector Search]
    BM25 & SEM --> HYBRID[Hybrid Fusion]
    HYBRID --> RERANK[Cross-Encoder\nReranker]
    RERANK --> CONTEXT[Top-K Context\nChunks]
end

subgraph GATEWAY_RAG["AI Gateway"]
    CONTEXT --> PROMPT[Prompt Assembly\n+ PII Redaction]
    PROMPT --> MLLM[Multi-LLM Router\nforge-router]
    MLLM --> CLAUDE_R[Claude]
    MLLM --> GEMINI_R[Gemini]
    MLLM --> GROQ_R[Groq]
    MLLM --> OAI_R[OpenAI]
end

subgraph OBS["Observability"]
    LANGFUSE[Langfuse\nLLM Tracing]
    GRAFANA_R[Grafana\nInfra Metrics]
    JUDGE[LLM-as-a-Judge\nQuality Eval]
end

MLLM --> LANGFUSE
CLAUDE_R & GEMINI_R & GROQ_R --> GRAFANA_R
LANGFUSE --> JUDGE

```

DevSecOps Pipeline — Security at Every Layer

```

flowchart LR
    DEV[Developer\nPushes Code] --> PRE

    subgraph PRE["Pre-Commit"]
        P1[Secret Scan]
        P2[Linting]
        P1 --> P2
    end

    PRE --> CI

    subgraph CI["CI Pipeline – Automated Gate"]
        C1[SAST\nCode Vulnerabilities]
    end

```

```

    C2[SCA\nDependency CVE]
    C3[Container\nImage Scan]
    C4[IaC Policy\nValidation]
    C5[Coverage\nThreshold]
    C1 --> C2 --> C3 --> C4 --> C5
end

CI -->|All Gates Pass| DEPLOY

subgraph DEPLOY["Pre-Deployment"]
    D1[Vault Secrets\nInjection]
    D2[K8s RBAC\nEnforcement]
    D3[Approval Gate\nfor Production]
end

DEPLOY --> RUNTIME

subgraph RUNTIME["Runtime – GitOps Provenance"]
    R1[Audit Trail\nGit = Source of Truth]
    R2[Anomaly Detection\nObservability Telemetry]
    R3[AI Gateway\nPII Redaction + Guardrails]
end

style CI fill:#1a1a2e,color:#eee
style DEPLOY fill:#16213e,color:#eee
style RUNTIME fill:#0f3460,color:#eee

```

Enterprise Observability Transformation @ Devo

```

flowchart TB
    subgraph BEFORE["Before – Siloed Monitoring"]
        B1[Team A Metrics\nCustom Dashboards]
        B2[Team B Alerts\nEmail-based]
        B3[No Telemetry\nStandards]
        B4[Manual Alert\nConfiguration]
    end

    subgraph AFTER["After – Unified Observability Platform"]
        subgraph COLLECT["Collection Layer"]
            OT1[OpenTelemetry\nDaemonSet]
            OT2[OpenTelemetry\nSidecar]
            PROM[Prometheus\nScraping]
        end

        subgraph STORE_OBS["Storage & Analysis"]
            LOKI2[Loki\nLog Aggregation]
            PROM_DB[Prometheus\nTime Series DB]
            ELK2[ELK Stack\nLog Analytics]
        end

        subgraph VIZ["Visualization & Alerting"]
            GRAF[Grafana\nDashboards]
            AM[Alertmanager\nRouting + Escalation]
        end
    end

```

```

GIT_ALERTS[GitOps\nAlert Config as YAML]
end

OT1 & OT2 & PROM --> LOKI2 & PROM_DB & ELK2
LOKI2 & PROM_DB --> GRAF
PROM_DB --> AM
GIT_ALERTS -->|Declarative Config| AM
GRAF --> SLACK2[Slack\nAuto-Notif]
end

```

```

BEFORE -->|Transformation| AFTER

```

Education

Degree	Institution	Year
MBA — Information Technology & Project Management	Sikkim Manipal University (Distance Education)	2012
B.Tech — Electronics & Communication Engineering	UPTU Lucknow	2009

Certifications

Certification	Issuer	Status
AWS Certified Solutions Architect — Associate	Amazon Web Services	Active (2023–2026)
AWS Certified Cloud Practitioner	Amazon Web Services	Active
ITIL Foundation Certification	Axelos	Active
ServiceNow Fundamentals	ServiceNow	Active
Claude 101 — AI Operational Certification	Anthropic	Active
AWS Media Services — Elemental	Amazon Web Services	Completed

Open Source & Portfolio

Project	Description	Links
AI FORGE	Enterprise AI Agent Operating System — Claude Code + AWS Bedrock	GitHub
forge-router	Multi-LLM AI Gateway — 8 providers, priority-sorted fallback chain	GitHub
Vikash Portfolio	Production Next.js portfolio — TypeScript + Tailwind CSS + Framer Motion	Live
Medium Articles	Technical writing on DevOps, AI, Platform Engineering	Medium

Current Learning Focus

```

quadrantChart
  title Learning Investment Matrix (2026)
  x-axis Low Priority --> High Priority
  y-axis Beginner --> Advanced

  quadrant-1 Invest Deeply

```

quadrant-2 Build Proficiency
quadrant-3 Explore
quadrant-4 Maintain

LangGraph: [0.85, 0.65]
MCP Protocol: [0.90, 0.70]
LLMOps: [0.88, 0.75]
RAG Architectures: [0.82, 0.80]
AI Governance: [0.75, 0.60]
Qdrant: [0.70, 0.55]
TypeScript: [0.65, 0.50]
AI Gateway Design: [0.92, 0.85]
Langfuse: [0.78, 0.60]
Multi-Agent Systems: [0.88, 0.68]

Professional Value Proposition

"I don't build AI features — I build AI-powered operational systems. I don't write Terraform modules — I write Custom Terraform Providers. I don't manage Kubernetes clusters — I manage 19 clusters across 7 global regions. The pattern is consistent: take hard problems, build something permanent, measure the outcome."

What I Bring	Why It Matters
15+ years across 6 domains	Rare breadth without sacrificing depth — from RF engineering to LLMOps
Production GenAI (not theory)	AI FORGE built, deployed, running — not a demo project
Custom Terraform Provider dev	Rare skill most senior engineers never develop
Observability-first architecture	Built GitOps-managed telemetry from scratch — not just configured existing tools
Media + Cloud + AI convergence	Uniquely positioned for AI-powered media platforms and enterprise AI infrastructure
90-day transformation playbook	DORA-first, data-driven, deliver quick wins before setting long-term roadmap

Contact

Channel	Details
Email	vikashjaiswal.486@gmail.com
Phone	+91-8588800287
LinkedIn	linkedin.com/in/linkedin2vikashjaiswal
GitHub	github.com/vikas0486
Medium	medium.com/@vikash.jaiswal
Location	Noida / Gurgaon / Delhi NCR, India
Availability	15 Days Notice · Open to Remote / PAN India
